

Analysis of semaphore implementations

The following implementation is provided at <https://sites.cs.ucsb.edu/~rich/class/cs170/notes/Semaphores/index.html>:

```
void P(sema *s){
    pthread_mutex_lock(&s->lock);
    s->value--;
    while (s->value < 0){
        if (s->waiters < -1 * s->value){
            s->waiters++;
            pthread_cond_wait(&s->wait,&s->lock);
            s->waiters--;
        }else{
            break;
        }
    }
    pthread_mutex_unlock(&(s->lock));
    return;
}
```

```
void V(sema *s){
    pthread_mutex_lock(&s->lock);
    s->value++;
    if (s->value <= 0){
        pthread_cond_signal(&s->wait);
    }
    pthread_mutex_unlock(&(s->lock));
}
```

`s->waiters` is initialized to 0 and is non-negative because for every decrement operation there is a preceding increment operation. Thus, for any `s->value >= 0` `s->waiters < -1 * s->value` is false. The following simplified implementation of P will be considered.

```
void P(sema *s){
    pthread_mutex_lock(&s->lock);
    s->value--;
    while (s->waiters < -1 * s->value){
        s->waiters++;
        pthread_cond_wait(&s->wait,&s->lock);
        s->waiters--;
    }
    pthread_mutex_unlock(&(s->lock));
}
```

Wlog, let `s->value == -2` and `s->waiters == 2`. Thread A calls V, increments `s->value` to -1, and signals with `pthread_cond_signal`. Thread A continues running, calls P, decrements `s->value` to -2, and is not pushed onto the queue of waiting threads. Thread B is awakened by the signal from `pthread_cond_signal` called by thread A, reacquires mutex, and decrements `s->waiters` to 1. Because `s->value == -2` and `s->waiters == 1`, thread B does not exit the while loop, increments `s->waiters` to 2 and is pushed back onto the queue of waiting threads. Thus, thread A "received" its own signal, avoided waiting, and the awakened thread B was pushed back onto the waiting queue of threads.

The above behavior reduces the role of the scheduler and can lead to thread starvation. The behavior is avoided by guaranteeing the queuing of each thread that

calls P (sem_wait) when the value of a semaphore is ≤ 0 after mutex acquisition in P (sem_wait), as provided in The Little Book of Semaphores by Allen B. Downey (Version 2.2.1):

```
void sem_wait(Semaphore *semaphore){
    mutex_lock(semaphore->mutex);
    semaphore->value--;
    if (semaphore->value < 0){
        do{
            cond_wait(semaphore->cond, semaphore->mutex);
        }while (semaphore->wakeups < 1);
        semaphore->wakeups--;
    }
    mutex_unlock(semaphore->mutex);
}

void sem_signal(Semaphore *semaphore){
    mutex_lock(semaphore->mutex);
    semaphore->value++;
    if (semaphore->value <= 0){
        semaphore->wakeups++;
        cond_signal(semaphore->cond);
    }
    mutex_unlock(semaphore->mutex);
}
```

The do...while loop guarantees the queuing of each thread that calls sem_wait, when `semaphore->value` ≤ 0 after mutex acquisition in sem_wait. `semaphore->wakeups` provides the accounting for ensuring that a signal from `cond_signal` does not awaken more than one thread due to the specification of pthread.